


Author the 76-task golden benchmark, wire in the two data modalities, and label every task for routing

A Authoring & data modalities



B Schema

The 22-column task schema




C Distribution

Task taxonomy across the 76 tasks




Build a hermetic, reproducible environment and pin everything before any model runs

Environment freeze – everything pinned before Stage 3 schedules a single run




Take the cartesian product of models, harness configs, tasks and seeds; dispatch one isolated cell per run



4 Agent Execution

Synth (on Hermes) runs each cell: plan, act through sandboxed tools under guards, then route by confidence

One cell = one model × harness × task × seed




5

Capture & Normalize


Collect the three output streams, diff the system state, and store everything for replay

Nothing is graded raw – every stream is normalized, diffed and indexed first




Route each task by gradability to three graders that compare captured outputs against the golden references

Three independent grading signals per task



Roll the grader signals into four metric families; compute variance across the N seeds

Four metric families → one scorecard per cell




8

Benchmark & Decision

Aggregate into a weighted leaderboard, track drift, and gate the billable-deployment decision

From scorecards to a defensible deployment decision





Determinism-by-Design Substrate

The cross-cutting backbone that makes every stage above reproducible, auditable and trustworthy

Eight guarantees, each anchored to the stage that enforces it

